



ATLAS Paper Draft

TRIG-2021-01

Version 3

Target journal: JINST

Comments are due by: August 18 2025

Supporting internal notes

Support Note: <https://cds.cern.ch/record/2777971/>

Implementation and performance of the Neural-Ringer electron selection algorithm in the high-level trigger of the ATLAS experiment

Electron trigger identification is a crucial input to many ATLAS physics analyses. To cope with the increase in peak LHC luminosity from 2017 to 2018, which raised the number of proton–proton interactions per beam-crossing to about 60, trigger algorithms were optimized to keep the background trigger rates low while retaining high efficiency for physics analyses. In the second half of 2017, the Neural-Ringer algorithm, based on a set of multi-layer perceptron models, replaced the cut-based selection algorithm used in the FastCalo decision step of high-level trigger sequence for events containing electrons above 15 GeV. This was done as part of the ATLAS trigger CPU optimization reduction campaign where it was the first time a neural-network method was employed as a baseline algorithm for event selection in the ATLAS trigger system. The selection algorithm resulted in an early rejection of at least 13 times more background than the previous algorithm strategy, while leaving the electron selection efficiency nearly unchanged.

Analysis Team

[*email:* atlas-TRIG-2021-01-analysis-team@cern.ch]

Jose Seixas, Denis Oliveira Damazio, Chris Meyer, Werner Spolidoro Freund, Joao Victor Da Fonseca Pinto, Micael Verissimo De Araujo

Editorial Board

[*email:* atlas-TRIG-2021-01-edboard-conveners@cern.ch]

Bertrand, Laforge (Paris LPNHE)
Fernando, Monticelli (La Plata) (chair)

Implementation and performance of the Neural-Ringer electron selection algorithm in the high-level trigger of the ATLAS experiment

The ATLAS Collaboration

Electron trigger identification is a crucial input to many ATLAS physics analyses. To cope with the increase in peak LHC luminosity from 2017 to 2018, which raised the number of proton–proton interactions per beam-crossing to about 60, trigger algorithms were optimized to keep the background trigger rates low while retaining high efficiency for physics analyses. In the second half of 2017, the Neural-Ringer algorithm, based on a set of multi-layer perceptron models, replaced the cut-based selection algorithm used in the FastCalo decision step of high-level trigger sequence for events containing electrons above 15 GeV. This was done as part of the ATLAS trigger CPU optimization reduction campaign where it was the first time a neural-network method was employed as a baseline algorithm for event selection in the ATLAS trigger system. The selection algorithm resulted in an early rejection of at least 13 times more background than the previous algorithm strategy, while leaving the electron selection efficiency nearly unchanged.

Contents

1	Introduction	3
2	Detector and Data acquisition	3
2.1	The ATLAS detector	3
2.2	The trigger system	4
3	Electron Identification	5
3.1	Offline object reconstruction and identification	5
3.2	Online reconstruction and identification by the Trigger	5
4	The Neural-Ringer algorithm	7
4.1	Building the rings	7
4.2	Motivation for using a set of neural networks	9
4.3	Data samples and training	9
4.4	Training and operation	10
5	Online deployment and performance	12
5.1	Measurement of the electron trigger efficiency	13
5.2	Assessment and deployment of Neural Ringer trigger Algorithm during Run 2 operation	14
5.3	Impact on CPU demands	15
6	Neural-Ringer vs cut-based triggers	17
6.1	Quadrant Analysis	19
6.1.1	Results	19
6.2	Homogeneity Tests	22
6.2.1	Results	22
7	Conclusion	24

1 Introduction

The ATLAS physics programme relies on an efficient trigger system [1] to record a highly signal-enriched subset of collision events produced by the Large Hadron Collider (LHC) at CERN. Electron triggers [2] are important for searches for new particles[3, 4], measurements of Standard Model cross-sections, and precise measurements of the properties of fundamental particles such as the Higgs [5, 6], Z and W bosons.

This paper describes the performance improvements achieved by the trigger system due to the introduction of a neural network based algorithm, the Neural-Ringer algorithm, to identify electrons. The Neural-Ringer became the baseline algorithm for selecting isolated electrons online by the ATLAS trigger system in the 2017 proton–proton collision at $\sqrt{s} = 13$ TeV in LHC Run 2.

The ATLAS detector and the trigger system are described in section 2. Electron identification both offline and online is described in section 3. Full details of the Neural-Ringer electron identification, including the training strategies are given in section 4. Online performance of Neural-Ringer and its comparison with respect to previous baseline algorithm are presented in section 5. The statistical analysis on the bias of the Neural-Ringer electron identification is described in section 6. Finally, the conclusions and prospects are presented in section 7.

2 Detector and Data acquisition

2.1 The ATLAS detector

ATLAS [7–9] is a multipurpose detector designed to observe particles produced in high-energy proton–proton (pp) and heavy-ion (HI) collisions. It is composed of a tracking detector in the innermost region around the interaction point, surrounded by calorimeters and muon chambers. The inner tracking detector (ID) is immersed in a 2 T magnetic field produced by a thin superconducting solenoid, and provides precise reconstruction of charged-particle tracks in a pseudorapidity¹ range $|\eta| < 2.47$. The innermost part consists of a silicon pixel detector. A silicon microstrip tracker surrounds the pixel detector with typically four layers of sensor modules. The outermost region of the tracker is covered by a transition radiation tracker (TRT). It consists of straw drift tubes filled with a gas mixture of 70% Xe, 27% CO₂ and 3% O₂ interleaved with polypropylene/polyethylene transition radiators creating transition radiation for particles with a large Lorentz factor.

Surrounding the inner tracking system, the ATLAS detector features a calorimeter system composed of both electromagnetic (EM) and hadronic (HAD) sections. This system is designed to ensure complete ϕ coverage and spans the pseudorapidity range $|\eta| < 4.9$, with enhanced granularity in the region corresponding to the inner detector coverage.

The EM calorimeter is a lead/liquid-argon (LAr) sampling calorimeter with an accordion-shaped geometry. It is segmented into two half-barrels ($-1.475 < \eta < 0$ and $0 < \eta < 1.475$) and two endcap modules ($1.375 < |\eta| < 3.2$). The transition area between the barrel and endcap sections ($1.37 < |\eta| < 1.52$)

¹ ATLAS uses a right-handed coordinate system with its origin at the nominal interaction point (IP) in the centre of the detector and the z -axis along the beam-pipe. The x -axis points from the IP to the centre of the LHC ring, and the y -axis points upward. Cylindrical coordinates (r, ϕ) are used in the transverse plane, ϕ being the azimuthal angle around the beam-pipe. The pseudorapidity is defined in terms of the polar angle θ as $\eta = -\ln \tan(\theta/2)$. The angular distance ΔR is defined as $\Delta R = \sqrt{(\Delta\eta)^2 + (\Delta\phi)^2}$. Transverse momenta and energies are defined as $p_T = p \sin \theta$ and $E_T = E \sin \theta$, respectively.

includes a considerable amount of inactive material [7]. Within the region where electrons are measured with precision ($|\eta| < 2.5$, excluding the transition region), the EM calorimeter is structured into three longitudinal layers based on shower depth. Each layer is composed by a set of calorimeter cells with a rectangular shape in the η, ϕ plane. The first layer features finely segmented cells in the η direction, which allow effective discrimination between isolated photons and closely spaced photon pairs resulting from $\pi^0 \rightarrow \gamma\gamma$ decays. For high transverse energy electrons and photons, the majority of energy is deposited in the second layer, with a cell granularity of 0.025×0.025 in (η, ϕ) space. The third layer captures energy deposits in the shower's tails. A thin presampler layer, located in front of the accordion calorimeter and covering $|\eta| < 1.8$, serves to correct for energy lost before reaching the calorimeter.

Beyond the EM calorimeter, three layers of hadronic calorimeters provide additional measurement capabilities. Additional background rejection for electrons and photons is achieved by detecting associated hadronic energy. The barrel hadronic calorimeter, covering $|\eta| < 1.7$, is a sampling calorimeter using iron and scintillator tiles, with light collection via wavelength-shifting fibers. The hadronic endcaps use copper/LAr calorimeters. In the forward region (FCal), copper–tungsten/LAr calorimeters are employed for both electromagnetic and hadronic measurements extending up to $|\eta| = 4.9$. All LAr-based calorimeters are enclosed within one barrel and two endcap cryostats.

The outermost component of ATLAS is the muon spectrometer (MS), which operates in the pseudorapidity range $|\eta| < 2.7$ and consists of three large toroidal magnet systems, each comprising eight coils. The magnetic field integral of the toroid system ranges from 2.0 to 6.0 Tm across most of its coverage. The MS integrates precision tracking chambers alongside fast-response trigger detectors.

2.2 The trigger system

A hierarchical two-level trigger system [1] is used to select events of interest produced in LHC beam collisions. The first-level (L1) trigger, implemented with custom-made electronics, utilizes the calorimeters at coarser granularity information to reduce the event rate from the 40 MHz bunch-crossing rate to below 100 kHz, of which about 35% corresponds to events with high-energy electromagnetic showers, potentially coming from electrons or photons. The L1 trigger also defines regions-of-interest (RoIs), situated around local regions with high transverse energy.

The RoIs defined by the L1 trigger are processed by the high-level trigger (HLT) implemented in software. The HLT operates in a large computing farm and reduces the rate of events written to disk to an average of about 1 kHz. The HLT uses fine-granularity calorimeter information, precision measurements from the muon spectrometer and tracking information from the ID, which are not available at L1. HLT reconstruction can be executed either within the RoIs identified at L1 or for the full detector (full-scan). The selection of particle candidates by the HLT is performed in a sequence of steps. First a set of fast reconstruction is performed, using only detector information within the RoI and faster and simpler algorithms to achieve early background rejection. Secondly a set of offline-like precision reconstruction algorithms are executed at a lower rate, having access to the data of the whole detector and performing the final decision for the Trigger.

A sequence of L1 and HLT trigger algorithms is called a ‘trigger’ and is meant to identify one or more particles of a given type and a given threshold of transverse energy or momentum. Electron triggers are meant to select events with one or more electrons in the final state identified by the detector. The ‘trigger menu’ defines the configuration of the trigger setting a full list of the L1 and HLT triggers and their thresholds. Menu composition and trigger thresholds are optimised for the LHC running conditions (beam

type, luminosity, etc.) to fit within the event acceptance rate and the bandwidth constraints of the data acquisition system of the ATLAS detector as well as the offline storage constraints.

3 Electron Identification

3.1 Offline object reconstruction and identification

The offline reconstruction of electrons and photons [10] is based on dynamic, variable-sized clusters of energy deposits, known as topo-clusters, formed from topologically connected electromagnetic (EM) and hadronic calorimeter cells [11]. These clusters help recover energy lost due to bremsstrahlung photons or electrons originating from photon conversions. After applying initial corrections for position and energy calibration, the topo-clusters are matched to tracks in the inner detector (ID) that have been re-fitted to account for bremsstrahlung effects, following the method outlined in Ref. [12], to reconstruct electron candidates. Topo-clusters that are either unmatched to any track or matched to a conversion vertex are reconstructed as photon candidates. Once candidates are formed, their energies undergo further recalibration for use in analyses.

Prompt electrons traversing the central region of the detector ($|\eta| < 2.47$) are identified using a likelihood-based (LH) method, which leverages several discriminating variables: the longitudinal and lateral shapes of EM showers, the quality of the track, matching between track and cluster, and particle identification information from the TRT. The probability density functions (pdfs) used in the LH approach are derived from $J/\psi \rightarrow ee$ decays for the transverse energy (E_T) range of 4.5–15 GeV, and from $Z \rightarrow ee$ decays for $E_T > 15$ GeV, as detailed in Ref. [10]. Separate pdfs are defined for each discriminant variable across different bins of E_T and η of the electron-candidate.

Three identification operating points are defined, each corresponding to progressively tighter requirements on the LH discriminant. The average efficiencies for electrons from electroweak processes (evaluated over the 2015–2017 dataset) are approximately: ‘loose’ (93%), ‘medium’ (88%), and ‘tight’ (80%).

3.2 Online reconstruction and identification by the Trigger

Electron reconstruction at the HLT stage is performed on each EM RoI provided by L1, which satisfies E_T and isolation requirements as specified by the trigger configuration. It proceeds in a series of sequential steps as shown in the schematic diagram in Figure 1. Lozenges represent feature-extracting algorithms which are responsible for building electron properties from the detector data (calorimeter shower shapes, transverse momentum, energy, position, etc). Rectangles represent hypothesis-testing algorithms, which apply selection requirements on those reconstructed properties.

Electron identification at the HLT selection starts with a fast calorimeter-reconstruction (FastCalo) step, which has two implementations: the original/baseline cut-based selection and the new neural-network (NN) approach using the Neural-Ringer algorithm presented in this paper. The latter implemented for triggering electrons with $E_T \leq 15$ GeV. FastCalo is followed by a fast tracking-reconstruction step, which build tracks for electron candidates only inside the RoI. A fast electron hypothesis algorithm applies requirements combining the results of both fast algorithms. Precision reconstruction algorithms are executed for trigger electron candidates satisfying the fast electron hypothesis selection.




Figure 1: This figure presents the sequence of algorithms used to select electron candidates at the High Level Trigger stage. Differences between the original baseline (no-ringer) and new (ringer) electron triggers lie in the first step, called FastCalo. Grey blocks represent algorithms in the fast stage, and red blocks represent precision-step algorithms.

In the precision calorimeter reconstruction step, the cluster reconstruction and calibration are similar to those for photons, and to those used offline in early Run 2 analyses [13]. Then, precision tracks within the RoI are extrapolated to the second layer of the EM calorimeter and are required to match the clusters within $|\Delta\eta(\text{track}, \text{cluster})| < 0.05$ and $|\Delta\phi(\text{track}, \text{cluster})| < 0.05$ rad. The precision electron selection combines the results of the precision calorimeter and precision tracking steps to build a likelihood (LH) discriminant with four operating points: ‘lhv loose’, ‘lh loose’, ‘lh medium’, and ‘lhtight’. The identification in the trigger is designed to be as close as possible to the offline version, except for the following differences: the discriminating variables used online have different resolutions; the momentum loss due to bremsstrahlung, p/p , is not accounted for in the online LH; the online selection uses $\langle\mu\rangle$ to estimate pileup instead the number of primary vertices used offline.

Triggers with ‘nod0’ suffixed to their names do not include the transverse impact parameter relative to the beam-line, d_0 , and its significance, in the online LH version. An additional, optional requirement of isolation denoted ‘ivar loose’ is also available for electron triggers. This tracking-only isolation is required to satisfy $p_T^{iso}(\Delta R^{var} < 0.2)/p_T < 0.10$ and is calculated similarly to the offline isolation operating

170 points.

171 4 The Neural-Ringer algorithm

172 The identification of electromagnetic objects at the FastCalo stage of the HLT must efficiently reject the
 173 large background from hadronic jets while operating under tight latency constraints. The Neural-Ringer
 174 algorithm was developed to enhance this discrimination capability by exploiting the approximately conical
 175 symmetry of electromagnetic showers in the calorimeter. Electrons and photons deposit most of their
 176 energy in a compact, centrally concentrated region, whereas hadronic jets exhibit wider and more irregular
 177 energy patterns extending into the hadronic calorimeter.

178 The idea behind the Neural-Ringer algorithm is to draw concentric rings perpendicular to the particle axis,
 179 estimate the energy deposited along each ring. This way the algorithm summarizes the shower profile
 180 as a function of the radial distance from the shower axis, converting the spatially distributed calorimeter
 181 information into a compact, physics-motivated representation suitable for neural-network classification.

182 4.1 Building the rings

183 The implementation of the Neural-Ringer define rings at every layer of both EM and HAD calorimeters.
 184 In the ATLAS calorimeter, the rings are built with a rectangular shape to match the grid layout of
 185 the calorimeter cells. Since this algorithm was designed to be used in the FastCalo step of the Trigger
 186 reconstruction, all the rings are confined within the area defined by the RoI and centered by the cell with
 187 highest energy deposited closest to the L1 position. As a general rule, a ring of size n in a given layer of the
 188 calorimeter consists of the cells lying in the border of an $n \times n$ cells area, cropped by the size of the RoI.




Figure 2: Sketch illustrating the ring-shaped energy description. The most energetic cell is in red, while the consecutive neighbouring rings are represented by alternating grey and black cells. Extracted from [2].

The actual definition of the rings is depicted in figure 2. In the EM calorimeter, rings are centred around the most energetic cell at each layer, while in the hadronic calorimeter the position of the most energetic cell in the second layer of the EM calorimeter is used as an axis. Formally, the cells that belong to a given ring of size n on a given layer l is given by equation 1 where $\eta_{i,l}$, $\phi_{i,l}$ represent the position of the cell i at detector layer l , the subindex “hot” refers to the cell with maximum energy in the RoI for the given layer; $h_{\eta,l}$ and $h_{\phi,l}$ are the sizes in η and ϕ of the cell at the layer l ; $\Theta_{\text{RoI},l}$ is the set of cells in the l^{th} layer which are within the RoI; and $\lfloor \cdot \rfloor$ is the *round* function. Table 1 shows the number of rings computed at each layer and the respective granularity.

$$R_{n,l} = \left\{ c_{n,l} \mid n = \left\lfloor \max \left(\frac{|\eta_{i,l} - \eta_{\text{hot},l}|}{h_{\eta,l}}, \frac{|\phi_{i,l} - \phi_{\text{hot},l}|}{h_{\phi,l}} \right) \right\rfloor, \forall c_{i,l} \in \Theta_{\text{RoI},l} \right\}, \quad (1)$$

Table 1: Nomenclature defining the Neural-Ringer algorithm layers and sections, as well as the respective parameters employed and the calorimeter sampling layers from which the cells were extracted. The Neural-Ringer algorithm parameters are dependent on the layer l but independent of η and E_T . The reconstruction parameters are the cell layer size in η ($h_{\eta,l}$), ϕ ($h_{\phi,l}$) and the number of rings to be computed in each layer (N_l).

Ringer		Calorimeter Sampling			Parameters		
Section	Layer (l)	Barrel	ITC	End-cap	$h_{\eta,l}$	$h_{\phi,l}$	N_l
EM	PS	PreSamplerB		PreSamplerE	0.025	0.1	8
	EM1	EMB1		EMEC1	0.003	0.1	64
	EM2	EMB2		EMEC2	0.025	0.025	8
	EM3	EMB3		EMEC3	0.050	0.025	8
HAD	HAD1	TileBar0 TileExt0	TileGap3	HEC0	0.1	0.1	4
	HAD2	TileBar1 TileExt1	TileGap1	HEC1 HEC2	0.1	0.1	4
	HAD3	TileBar2 TileExt2	TileGap2	HEC3	0.2	0.1	4

For every layer of the EM and HAD calorimeters, the transverse energy deposited in the cells lying along each ring is summed. Across all layers, a total of 100 ring transverse energies are computed, normalized to the RoI E_T , ordered from the innermost EM region to the outermost hadronic layer and stored into a single vector. Each index in that vector identifies a ring. The energy normalization of the rings is performed as in Eq. (2) where r_k or r_i is the sum of transverse energy of the cells in the ring with index k or i , and r'_k is the normalized ring energy value with index k . The absolute value is used because of negative noise accumulation from some cells within the reconstruction window that were not excited by the particle shower. This procedure was initially proposed and examined in Ref. [14] as a way of preserving the shower energy profile as fractions of the total energy.

This procedure compresses the information of thousands of calorimeter cells per RoI into a vector of one hundred variables while preserving both the longitudinal and lateral structure of the energy deposition.

$$r'_k = \frac{r_k}{\sum_{i=1}^{100} r_i}, \quad \forall k \in \{1, \dots, 100\}, \quad (2)$$

Figure 3 shows the distribution of ring energy as a function of the ring number (or vector index) for electrons and QCD jets obtained in simulated events. Rings within vertical dashed lines belong to the same layer of the same calorimeter. The distribution of the ring energy for layer 2 of the EM calorimeter (EM2) is shown zoomed within the figure.




Figure 3: Simulated average ring profiles for jets and electrons in the bin $30 \leq E_T < 40$ GeV and $0.0 \leq \eta < 0.8$ for each layer of the calorimeter. The profile in the inset shows the EM1-layer ring energies on a logarithmic scale. Almost all the energy of the electrons is absorbed by the electromagnetic layers and no energy deposit is expected in the hadronic calorimeter. Instead, a significant fraction of the jet energy is absorbed by the hadronic layers and only a part of it is absorbed by the electromagnetic calorimeter. The shaded area indicates the standard deviation of each ring energy.

4.2 Motivation for using a set of neural networks

The detector response and discrimination power between electrons and QCD jets changes with the candidate E_T and also with its η position, given the different granularities of the calorimeters at different pseudorapidities. Therefore a set of multi-layer perceptrons (MLPs) trained on different regions of $E_T \times |\eta|$ is used to classify the normalized vector of ring energies. Splitting the training into multiple small models was found to have better discrimination power and to be more efficient in terms of computing resources with respect to a single training inclusive in E_T and $|\eta|$ regions.

Low-complexity models, with a single hidden layer and few neurons were chosen for the MPL set in order to cope with limited CPU resources at the Fast Processing step of the Trigger and to minimize the latency.

4.3 Data samples and training

The Neural-Ringer algorithm implemented online during 2017 was trained with simulated samples of $Z \rightarrow e^+e^-$ events for signal and QCD dijet events for background. The electrons from the signal sample are selected using tag-and-probe method described in [2], in which one electron (the tag) is used to trigger the event and its requested to satisfy tight identification and isolation requirements offline. A second electron (the probe) is required to be in the event on which almost no requirement is applied, except for the

invariant mass of the Tag and Probe system to be within the mass of the Z boson. This makes the probe almost unbiased with respect to offline identification and completely unbiased with respect of the trigger selection. To enrich the sample of electron backgrounds in the QCD simulation, only events where the sum of transverse energy of particles of the hard scatter (excluding muons and neutrinos) in a window of $\Delta\eta \times \Delta\phi = 0.1 \times 0.1$ exceeds 17 GeV are selected. Both $Z \rightarrow e^+e^-$ and hadronic jets from QCD hard scatter simulated events included pileup of 60 interactions per bunch crossing covering running conditions of full Run-2.

For 2018 operation, a purely data-driven strategy was used, using LHC pp collisions at $\sqrt{s} = 13$ TeV collected by the ATLAS detector during 2017 run after the application of beam, detector and data quality requirements [15].

4.4 Training and operation

A dedicated MLP model is defined for each region of the $E_T \times |\eta|$.

Threshold values for the discriminating variables were defined based on 2016 collision data, aiming for a signal efficiency comparable to that of the cut-based trigger. The training procedure and decision-making process were identical across all phase-space regions. For the Neural-Ringer, each MLP serves as an expert model corresponding to a specific phase-space region, with its own topology and parameters.

The model parameters were optimized using events selected from simulated datasets. The network architecture consists of a single fully connected hidden layer containing between 5 and 20 hidden units, an input layer with 100 nodes (one for each ring), and an output layer with a single neuron. The activation function for both the hidden and output neurons is the hyperbolic tangent, which is commonly used in MLP designs [16].

To assess the statistical fluctuations of the efficiency measurements with respect to the model parameters obtained from the training procedure, the stratified k -fold cross-validation method [16], with $k = 10$, was used, although repeating the training and testing procedures on different randomly chosen subsets increased the computational cost. The stratified k -fold is one of the most common cross-validation techniques and consists of partitioning the dataset into k disjoint subsets, each model being tested with the i^{th} subset after being trained with the other ones. The dataset stratification follows the empirical order of the event selection, in order to allow straightforward job parallelization, i.e. random permutation is not employed.

For each cross-validation sort and working point, 100 networks, initialized as in [17], are optimized by back-propagating the mean squared error through the RPROP algorithm [18]. This repeated model initialization seeks to avoid poor sub-optimal solutions due to the use of gradient-based algorithms. For each training procedure, only three models out of the 100 are retained: two for retrieving the best efficiency when fixing the sensitivity or false positive probabilities to the baseline FastCalo efficiency and another that gives the maximum SP value.

The maximum SP is the value that best relates the balance between signal detection probability and the probability of fake signal in the ROC curve (i.e., the knee of the curve) and is defined as:

$$SP = \sqrt{\sqrt{P_D(1 - F_A)} \cdot \frac{P_D + (1 - F_A)}{2}}$$

where P_D is the signal detection probability (or sensitivity) and F_A is the probability of a fake signal (false positive). An early stopping [16] technique is employed to avoid model over-training. The selection of the optimal iteration and operating model is based on a heuristic approach (multi-stop) [19].

Computational resources were used to tune 1.3 million shallow-learning neural networks, which resulted in a set of 20 MLPs for each one of the four working points used by electron triggers in the HLT. The training strategy for 2017 data-taking considered only four regions in $|\eta|$ ($0.0 \rightarrow 0.8 \rightarrow 1.37 \rightarrow 1.57 \rightarrow 2.47$) and five regions in E_T ($15 \text{ GeV} \rightarrow 20 \text{ GeV} \rightarrow 30 \text{ GeV} \rightarrow 40 \text{ GeV} \rightarrow 50 \text{ GeV} \rightarrow \infty$). Apart from few exceptions, the models in the set employ five neurons in the hidden layer.

After the training stage, decisions are made by applying a threshold requirement to the one-dimensional output node of each expert MLP. As in the HLT likelihood [20] algorithm, the threshold value used in the Neural-Ringer algorithm is computed as a linear function of $\langle\mu\rangle$, the mean number of pileup interactions per bunch crossing. However, non-linear behaviour was observed near the asymptote of the output activation function (see Figure 4(a)) when using the ‘medium’ or ‘tight’ operating point. It was observed that this behaviour can be avoided by replacing the hyperbolic tangent activation function in the output neuron with a linear function (Figure 4(b)) for operation after the training stage is completed. This strategy ensures linear behaviour of the targeted² signal efficiency with respect to the online pileup estimator, similar what was used in [20]. To account for Monte Carlo and collision data mis-modelling (see Figure 5), the parameters of the linear threshold correction were derived using 2016 collision data and an upper bound of $\langle\mu\rangle = 40$ pileup interactions (i.e. the corrected threshold value is constant for pileup values above the upper bound).


It was observed that the rings in the region $2.37 < |\eta| < 2.47$ had particular profiles due to the lack of EM1 cells in this region and required a specific discriminant threshold in order to keep the signal efficiencies close to those of the cut-based triggers. In order to avoid retraining all models, it was decided to split the last $|\eta|$ region in two, from $1.37 \rightarrow 2.47$ to $1.37 \rightarrow 2.37 \rightarrow 2.47$, and duplicate all models, for this region, to the new bin. Finally, the set operated in 2017, which combines every η boundary with each E_T boundary, has a total of 25 MLP models. The set’s boundaries are defined in Table 2.

Table 2: E_T and η boundaries used to create the set of neural networks.

Model Adjustment	
E_T [GeV] boundaries	$ \eta $ boundaries
$15 \leq E_T < 20$	$0.00 \leq \eta < 0.80$
$20 \leq E_T < 30$	$0.80 \leq \eta < 1.37$
$30 \leq E_T < 40$	$1.37 \leq \eta < 1.54$
$40 \leq E_T < 50$	$1.54 \leq \eta < 2.37$
$E_T \geq 50$	$2.37 \leq \eta < 2.47$

On the other hand, the 2018 training procedure employed a single set structure for all working points and used the 2017 collision data, selected with $Z \rightarrow e^+e^-$ tag-and-probe event selection, for training, as was the case for the derivation of the offline and final HLT likelihood models [13]. Most of the procedure was kept unchanged for 2018. The developments for 2018 data benefited from the data conditions being similar to those in 2017. To simplify the training strategy, the selection of three models for each training was

² same signal detection probability of the online precision electron LH algorithm, increased by a factor of 50%, for each one of the four trigger operation points.



(a) MLP output with a hyperbolic tangent activation function in the output neuron versus pileup.

(b) MLP output with a linear activation function in the output neuron versus pileup.

Figure 4: Neural-Ringer output as a function of pileup, $\langle\mu\rangle$, for $Z \rightarrow e^+e^-$ probes in 2016 collision data selected using the training criteria. The computed threshold per $\langle\mu\rangle$ unit for achieving the targeted efficiency is shown as a full blue circle. The black line is a linear fit to the thresholds. The output space is computed using the trained model without modifications in (a), while in (b) the activation function in the output neuron is replaced with a linear function.

abandoned, keeping only the maximum SP model, as the previous approach was more complex without a clear increase in trigger efficiency.

Likewise, MLPs with 5 to 10 units in the hidden layer were tuned, as most models did not require more than 10 units in 2017. The training optimizer was changed to *Adam* [21], instead of RPROP, and the MLP model topology was modified to have a hidden layer with fully connected neurons, with rectified linear units (RELU) as the activation function, and an output layer with one neuron using a sigmoid activation function. Specialized MLPs were implemented in the region $2.37 < |\eta| < 2.47$, where different ring profiles were observed before 2017 operation, resulting in the operation of 25 MLPs. Finally, the upper bound for threshold corrections was set to $\langle\mu\rangle = 100$ pileup interactions in order to extrapolate the operation of the models to a never-reached pileup level.

5 Online deployment and performance

The Neural-Ringer algorithm was implemented in four stages during 2017–2018 in LHC Run 2:

- i A development stage up to early 2017, where the Neural-Ringer's performance was estimated rerunning the trigger decision offline.
- ii The commissioning stage in early 2017 runs (5.4 fb^{-1}), where all primary triggers were duplicated and both the Neural-Ringer and cut-based algorithm operated in the FastCalo stage of the HLT.




Figure 5: Differences between simulation samples (filled histogram) and collision data (empty histogram) obtained in 2016 $Z \rightarrow e^+e^-$ tag-and-probe for one region of the phase space. The comparison is made between the normalized distributions of the output of the neural network (the discriminant) obtained for these two cases. The shift in discriminant value between simulation and collision data is caused by differences between the simulation and data rings used as input to the model.

- iii The operation stage, starting after Technical Stop 1 (TS1) in 2017. Triggers employing the Neural-Ringer were used as the main triggers, and a small set of cut-based triggers were kept for validation purpose. Final validation of the Neural-Ringer was performed by comparing the main triggers and the cut-based triggers from this period (39.0 fb^{-1}) as shown in Section 5.2 and Section 6.
- iv The Neural-Ringer-only stage, during 2018. Duplicated cut-based triggers were removed. The performance of the Neural-Ringer in 2018 is presented in Section 5.2.

This section reports the operation of Neural-Ringer in each of these periods. In addition, CPU-time usage is discussed in Section 5.3.

5.1 Measurement of the electron trigger efficiency

The electron trigger efficiency is measured using the tag-and-probe method, which provides a sample of electrons selected without bias from the trigger decision. This technique allows unbiased estimation of the trigger performance by identifying probe electrons in events where a tag electron satisfies stringent offline and trigger requirements.

The efficiency is defined as the fraction of probe electrons that are accepted by the trigger, among those that pass the full offline selection:

$$\epsilon_{\text{trigger}} = \frac{N_{\text{trigger} \cap \text{offline}}}{N_{\text{offline}}} \quad (3)$$

where $N_{\text{trigger} \cap \text{offline}}$ is the number of probe electrons that pass both the offline selection and are accepted by the trigger, and N_{offline} is the total number of probe electrons passing the offline selection.

5.2 Assessment and deployment of Neural Ringer trigger Algorithm during Run 2 operation

The operation of Neural-Ringer started after the technical stop period (TS1) occurred in the middle of 2017. For monitoring purposes, a single-electron trigger was used. This trigger requires $E_T > 28$ GeV, with the selection criteria ‘tight’ and with ‘loose’ isolation criteria.

Signal efficiencies for both cut-based and Neural-Ringer triggers, evaluated for tag-and-probe electron candidates from 2017 data, are compared in Figure 6. The results show no significant differences, on the order of a few per mille, between the two algorithms.

Although a slight efficiency loss is observed at higher $\langle\mu\rangle$ values in Figure 6(c) for the Neural-Ringer trigger, the overall electron efficiency was maintained nearly constant, consistent with the objective of the Neural-Ringer algorithm’s training. This strategy is further validated by the results.

The observed efficiency loss occurs after the application of the linear threshold correction at $\langle\mu\rangle = 40$ during 2017.

Other important triggers were assessed by comparing the trigger efficiencies in the 2017 collision data collected before and after switching to the Neural-Ringer algorithms, as shown in Figure 7. The set of triggers evaluated was three. One of these triggers requires an electron candidate with $E_T > 17$ GeV, seeded by the first-level trigger with $E_T > 15$ GeV, and satisfying the loosest identification criteria (lhvloose). On the other hand, another requires an electron candidate with $E_T > 26$ GeV satisfying ‘tight’ identification (lhtight) with ‘loose’ isolation (ivarloose), and the last requires an electron candidate with $E_T > 60$ GeV satisfying ‘medium’ identification (lhmedium). The electron efficiency was also kept nearly unchanged for these relevant single-electron triggers. As mentioned above, the results in Figure 7 are computed for two different data-taking periods: one with, and the other without, the Neural-Ringer algorithm.

The presence of duplicated triggers, one configured with the Neural-Ringer algorithm and the other cut-based, allowed a direct comparison of their rates for fake electron candidates in collision data. The overall fake-electron rate is reduced by a factor of 13.75 at the FastCalo step. It can be seen on the left side of Figure 8 that the improvement at the FastCalo step is similar for all regions in the evaluation with respect to E_T , η , and $\langle\mu\rangle$ variables, which is particularly interesting when considering the low E_T and endcap regions. Besides improving early fake-electron rejection, the use of the Neural-Ringer also resulted in a factor of two reduction in the final fake-electron rate from the HLT (see right side of Figure 8), mostly coming from the calorimeter’s barrel-endcap transition regions.

For the 2018 data-taking period, the Neural-Ringer was retrained using collision data, contributing to improved electron identification performance. The signal efficiency of the lowest-transverse-energy-threshold unrescaled trigger increased by approximately 1%, primarily due to several minor enhancements in the full electron selection applied both online and offline. In parallel, a modest reduction in fake-electron acceptance was observed at the FastCalo stage for certain unrescaled triggers relative to the 2017 post-TS1 configuration. Specifically, the lowest- E_T -threshold triggers ($E_T > 17, 26, 28$) achieved a reduction factor of 1.19, while the highest- E_T -threshold trigger ($E_T > 60$) reached a more significant reduction factor of 1.69.




Figure 6: The efficiency of single-isolated-electron triggers with the Neural-Ringer algorithm and with the cut-based algorithms as a function of the offline electron (a) E_T , (b) η , and (c) $\langle\mu\rangle$ requiring $E_T > 28$ GeV after its deployment in July 2027. Efficiency is given with respect to the offline Tight identification working point. For (b) and (c), only offline candidates with $E_T > 29$ GeV are used.

5.3 Impact on CPU demands

The overhead required to compute the NN decision was evaluated by running the single-isolated-electron trigger requiring $E_T > 17$ GeV trigger with the Neural-Ringer algorithm or cut-based approach, in two individual non-concurrent executions using the same dedicated node³ on the 2017 Enhanced bias [22] (EB) offline data. Hence, the measurements are performed under pileup conditions, with the reconstruction algorithm executing for multiple RoIs in the same bunch-crossing event.

³ This performance study was performed using a Xeon Phi 7120 (1.7 GHz, 32 threads) node with a memory of 256 Gb@1333 running a SL6 Operating System.




Figure 7: Efficiency of three single-electron triggers as a function of (a) E_T and (b) η for 2017 collision data. Open (filled) markers show the efficiency measurements in runs before (after) the deployment of the Neural-Ringer, thus referring to triggers executed without (with) the Neural-Ringer algorithm.


It should be mentioned that, to reduce implementation effort, electron triggers using Neural-Ringer information compute their decision after the reconstruction of cut-based features, but the cut-based decision is not computed. As a result of this, the Neural-Ringer electron triggers always demand additional FastCalo processing time. Figures 9 and 10 show comparisons of the total processing time per call for the feature extraction and hypothesis-testing algorithms in the FastCalo step, respectively (see the related FastCalo blocks in Figure 1).

A multi-modal structure can be observed in the feature extraction (see Figure 9) distributions of both trigger configurations and it is associated with the algorithm for retrieving the EM2 cell energies and building related shower-shape variables. The computation of the ring variables, which is the only difference between the triggers in the FastCalo feature extraction, requires an additional average processing time of 0.18 ± 0.03 ms/event.


A smaller contribution comes from hypothesis testing (see Figure 10), in which the ring variables are normalized, the discriminating variable's value is computed, compared with the selection requirement, and stored in a persistent file format. Hypothesis testing for the Neural-Ringer exhibits significantly higher computational demands than the cut-based selection, with processing times reaching up to 0.14 ms/event. This overhead is predominantly due to the storage of the neural-network output, which requires 0.07 ± 0.02 ms per event. Despite this, the hypothesis testing remains less time-intensive than the feature extraction stage. Additionally, the time increase is negligible in comparison with the average processing time of a single-electron trigger across the entire HLT (≈ 500 ms in Run 2).

Taking these values into account, the Neural-Ringer can require up to 50% more FastCalo CPU time per event than the trigger without the Neural-Ringer.⁴ Nonetheless, the Neural-Ringer provides a more


⁴ The CPU-time increase is expected to depend on the trigger configuration because of the presence of pileup. The reported results are for the single-electron trigger, requiring $E_T > 17$ GeV, with a very loose operation point.




(a)



(b)



(c)



(d)

Figure 8: Efficiency of selecting background at the FastCalo (left) and HLT (right) steps as a function of E_T (top) and η (bottom) with a single-isolated-electron trigger requiring $E_T > 28$ GeV and ‘tight’ selection, with and without the Neural-Ringer algorithm, employing 2017 collision data collected after the Neural-Ringer algorithm deployment.

discriminating selection, which reduces the number of CPU-demanding triggers by decreasing the need to process fake electrons in subsequent steps that are more computationally intensive.

6 Neural-Ringer vs cut-based triggers

This section presents a detailed comparison between the Neural-Ringer and the cut-based electron trigger algorithm used in the ATLAS experiment during the 2017 and 2018 data-taking periods. The study aims to identify differences in trigger decisions and to evaluate potential biases introduced by the Neural-Ringer within the electron trigger sequence.




Figure 9: Total processing time per call for the feature extraction algorithms in the FastCalo step of an electron trigger requiring $E_T > 17$ GeV and Loose identification criteria, with the Neural-Ringer (blue) or the cut-based approach (red), on EB events (peak $\langle \mu \rangle = 45$). The inset on the right shows the processing time per call for only the ring variable extraction.




Figure 10: Total processing time per call for the hypothesis-testing algorithms in the FastCalo step of an electron trigger requiring $E_T > 17$ GeV and Loose identification criteria, with the Neural-Ringer (blue) or the cut-based approach (red), on EB events (peak $\langle \mu \rangle = 45$). The centre histogram represents the processing time per call to store the whole Neural-Ringer output in persistent file format.

The comparison is performed using a dataset from 2017, selected via the $Z \rightarrow e^+e^-$ tag-and-probe method [2]. In each event, one electron—designated as the probe electron—was not involved in the trigger-level selection, enabling a direct comparison of algorithm performance. Only probe electrons that satisfy the offline electron identification criteria are considered. The initial analysis applies the

Neural-Ringer selection to the probe electron to assess trigger performance (Quadrant Analysis).

To further investigate algorithm-induced biases, a second dataset is used. This dataset was collected during a dedicated 2017 run in which both the cut-based and Neural-Ringer algorithms were simultaneously active. The corresponding tag-and-probe sample enables a direct evaluation of selection biases by comparing the trigger decisions applied to the tag electron (Homogeneity Tests).

6.1 Quadrant Analysis

The Quadrant Analysis was developed to compare the decisions of two classifiers. Given signal candidates and two classifiers, these are the possibilities: either both accept or reject the candidates, or only one accepts the candidates.

The four possible trigger decisions, referred to as quadrants, are used to categorize events. The shower-shape variables employed in the offline electron likelihood selection are then compared across these quadrants to identify any potential bias introduced by the trigger step. Furthermore, in this analysis, it is also possible to evaluate the disagreement between two classifiers. The disagreement between two classifiers **A** and **B** can be defined as:

$$\text{disagreement} = \frac{N_{\mathbf{A}} + N_{\mathbf{B}}}{N_{\mathbf{A}} + N_{\mathbf{B}} + N_{\mathbf{AB}} + N_{\neg\mathbf{AB}}} \quad (4)$$

where $N_{\mathbf{A}}$ ($N_{\mathbf{B}}$) is the total number of candidates accepted exclusively by classifier **A** (**B**) and $N_{\mathbf{AB}}$ ($N_{\neg\mathbf{AB}}$) is the total number of candidates accepted (rejected) by both classifiers.

In this approach, only probe electron candidates selected using the $Z \rightarrow e^+e^-$ tag-and-probe method in collision data are considered. The offline selection applied to these candidates corresponds to the tight working point used as the trigger classifier. The Results in Section 6.1.1 show the trigger selection performance using the tight requirement for duplicated triggers during 2017.

6.1.1 Results

The shower-shape variables (R_η , R_ϕ , E_{ratio} , and R_{had}) [2] used in the offline likelihood selection are compared across the four trigger decision categories (quadrants).

As shown in Figure 11, the disagreement—defined as cases where the classifiers make different decisions—is typically at the level of 1%, regardless of the shower-shape variable used in the offline electron identification likelihood. Here, only the results for $0.6 < |\eta| < 0.8$ and $35 < E_T < 40$ GeV are shown, but this behaviour is also observed for other regions of the $E_T \times |\eta|$ phase space. One should note that the working points of the two triggers are not exactly the same, although the Neural-Ringer aims to achieve the same signal efficiency as the cut-based strategy and was operating as desired. In other words, the count differences between the blue and red profiles in Figure 11 are mainly due to small efficiency differences between the two triggers.

The shower-shape variables exhibit similar distributions across the four quadrants, although slight shifts between the blue and red profiles are visible in Figure 11. This is the case for R_η and R_ϕ , where the Neural-Ringer trigger consistently selects slightly more events in the desired region of the histogram closer to one, i.e. ones with slightly narrower showers in η and ϕ , respectively, in the EM2 layer. The

Neural-Ringer trigger's behaviour shows an even smaller effect in R_{had} , where tighter tails are observed. The E_{ratio} variable also shows the Neural-Ringer trigger has a slight bias towards selecting more events in the tail region. Although similar behaviour is seen in other regions of the $E_{\text{T}} \times |\eta|$ phase space.




Figure 11: Quadrant Analysis plots for a selection of shower-shape variables used in the offline electron identification algorithm for the $0.6 < |\eta| < 0.8$ and $30 < E_T < 35 \text{ GeV}$ region. The top pad in each figure shows the raw number of candidates for the four mutually exclusive cases: accepted both with and without the Neural-Ringer (black); accepted only with the Neural-Ringer (blue); accepted only with the Cut-based (red); rejected both with Neural-Ringer and with Cut-based (grey). The bottom pad shows the disagreement as defined in Eq. (4). The absence of points of disagreement is due to the absence of events in the bins for the blue and red histogram cases.

6.2 Homogeneity Tests

This analysis evaluates the impact of using different trigger strategies—Neural-Ringer-based and cut-based—on the offline likelihood identification of tag electrons. To perform the study, a pair of duplicated single-electron triggers were used online. Both required $E_T > 28$ GeV, applied the ‘tight’ identification criteria, and included ‘loose’ isolation requirements. One version incorporated the Neural-Ringer-based selection, while the other did not. The same offline identification criteria used in the likelihood optimization described in [13] were applied throughout the analysis.

The problem is addressed by applying homogeneity tests to histograms [23] to identify potential differences in the distributions of physics quantities resulting from the use of different trigger configurations. These tests are based on a test originally developed by Pearson [24] and popularly employed in many fields beyond high-energy physics, e.g. social sciences [25] and health [26], where they are usually called contingency or consistency tests.

To enable homogeneity tests without adapting them to ATLAS-specific setups, the 2017 collision data were split into two statistically independent groups. These groups were constructed by alternating short time intervals to ensure similar data-taking conditions.

The p -value represents the probability of observing a result at least as extreme as the one measured, assuming the null hypothesis is true. In this case, the null hypothesis states that the distributions of a given variable are the same in both groups. Comparing p -values between the groups helps assess whether the trigger configuration affects the shape of physics distributions.

To further investigate possible distortions introduced by the trigger algorithms, a bin-by-bin comparison is performed on histograms of shower-shape variables used in the offline electron identification. The signed χ -value, defined in Equation 5, is used for this purpose. In this expression, $r_{i,j}$ and $b_{i,j}$ are the event counts in the j^{th} bin and i^{th} $E_T \times |\eta|$ region, collected by the group using the trigger with and without the Neural-Ringer, respectively. This equation quantifies the deviation between the two trigger configurations and allows both positive and negative contributions to be visualized.

$$\chi_{i,j}^s = \frac{(r_{i,j} - b_{i,j})}{\sqrt{b_{i,j}}}, \quad (5)$$

6.2.1 Results

Regardless the trigger configuration, the homogeneity tests between the two groups result in similar p -values. The comparison of shower-shape distributions between the two statistically independent groups reveals fluctuations consistent with statistical uncertainty. No significant deviations are observed that would indicate systematic differences arising from the trigger configuration.

Figure 12 shows the residuals (bottom) for the R_η , R_ϕ , E_{ratio} , and R_{had} calorimetry variables when the reference histogram is obtained from collision data collected by the trigger without the Neural-Ringer for the first group (top). The residuals are dominated by statistical fluctuations (i.e., the residual values fluctuate around a central value), and are within the expectations from the homogeneity hypothesis for most variables and regions. Similar behavior is observed for track variables.




Figure 12: **Top:** Profile histograms of the R_η , R_ϕ , E_{ratio} , and R_{had} calorimetry variables used in the offline likelihood selection, shown for the $E_T > 40$ GeV and $0.0 < |\eta| < 0.8$ region. The distributions are obtained from two statistically independent data groups: one collected using the trigger without the Neural-Ringer (blue area histogram), and the other using the trigger with the Neural-Ringer (black line histogram). The assignment of trigger type to group is arbitrary and alternated across time intervals to ensure statistical independence. **Bottom:** Bin-by-bin residuals computed using the signed χ^s statistic (black circles, see Eq. 5), which quantify the deviation between the two trigger configurations. The horizontal line at zero (purple dots) serves as a visual reference indicating no deviation.

7 Conclusion

The Neural-Ringer has an innovative design for electron selection based on calorimetry information. With the goal of satisfying specific trigger system demands, feature extraction from concentric ring energy sums provides an efficient description, which is quickly derived through simple operations. Neural-network models are employed to exploit non-linear correlations between rings, and the optimization of this strategy culminated in the deployment of a set of models specific to each region of pseudorapidity and transverse energy.

In the second half of 2017, the Neural-Ringer algorithm replaced a cut-based strategy in the FastCalo decision step of the HLT and became the baseline algorithm for triggering events containing electrons with E_T above 15 GeV, as part of the ATLAS online system's CPU-time reduction campaign. Furthermore, it was the first time a neural-network method was employed as a baseline event-selection algorithm in the ATLAS trigger system.

During 2017 operation, the Neural-Ringer algorithm trained on simulated data reduced fake-electron trigger rates by a factor of 13 in the FastCalo step and by a factor of 2 overall in the HLT with respect to the offline likelihood algorithm. At the same time, the Neural-Ringer achieving high electron-selection efficiencies very similar to those of relevant single-electron triggers without the Neural-Ringer. For 2018, the Neural-Ringer was trained with collision data, resulting in an estimated additional fake-rate reduction factor of 1.2–1.7 for fake-electron triggers in the FastCalo step with respect to 2017.

The Neural-Ringer was also studied using offline likelihood as a reference to obtain additional insights into its behavior by using two equivalent single-electron triggers with and without the Neural-Ringer. The ‘quadrant analysis’ showed a high level of agreement between the two triggers, with the Neural-Ringer trigger showing a slight bias towards the signal region in the disagreement profiles of some calorimetry-based variables. Finally, ‘homogeneity tests’ were performed on shower-shape and tracking variables used in the offline likelihood identification. The comparison between data collected with the Neural-Ringer-based and cut-based triggers yielded similar p-values, indicating no observable bias in these variables introduced by the trigger configuration. This confirms that the Neural-Ringer trigger does not distort the inputs to the offline selection. As a result, the improved efficiency–rejection performance observed with the Neural-Ringer trigger can be attributed to its online behavior, without compromising the integrity of the offline identification

Acknowledgements

We thank CERN for the very successful operation of the LHC and its injectors, as well as the support staff at CERN and at our institutions worldwide without whom ATLAS could not be operated efficiently.

The crucial computing support from all WLCG partners is acknowledged gratefully, in particular from CERN, the ATLAS Tier-1 facilities at TRIUMF/SFU (Canada), NDGF (Denmark, Norway, Sweden), CC-IN2P3 (France), KIT/GridKA (Germany), INFN-CNAF (Italy), NL-T1 (Netherlands), PIC (Spain), RAL (UK) and BNL (USA), the Tier-2 facilities worldwide and large non-WLCG resource providers. Major contributors of computing resources are listed in Ref. [27].

We gratefully acknowledge the support of ANPCyT, Argentina; YerPhI, Armenia; ARC, Australia; BMWFW and FWF, Austria; ANAS, Azerbaijan; CNPq and FAPESP, Brazil; NSERC, NRC and CFI,

Canada; CERN; ANID, Chile; CAS, MOST and NSFC, China; Minciencias, Colombia; MEYS CR, Czech Republic; DNRF and DNSRC, Denmark; IN2P3-CNRS and CEA-DRF/IRFU, France; SRNSFG, Georgia; BMBF, HGF and MPG, Germany; GSRI, Greece; RGC and Hong Kong SAR, China; ISF and Benoziyo Center, Israel; INFN, Italy; MEXT and JSPS, Japan; CNRST, Morocco; NWO, Netherlands; RCN, Norway; MNiSW, Poland; FCT, Portugal; MNE/IFA, Romania; MESTD, Serbia; MSSR, Slovakia; ARRS and MIZŠ, Slovenia; DSI/NRF, South Africa; MICINN, Spain; SRC and Wallenberg Foundation, Sweden; SERI, SNSF and Cantons of Bern and Geneva, Switzerland; MOST, Taipei; TENMAK, Türkiye; STFC, United Kingdom; DOE and NSF, United States of America.

Individual groups and members have received support from BCKDF, CANARIE, CRC and DRAC, Canada; CERN-CZ, FORTE and PRIMUS, Czech Republic; COST, ERC, ERDF, Horizon 2020, ICSC-NextGenerationEU and Marie Skłodowska-Curie Actions, European Union; Investissements d'Avenir Labex, Investissements d'Avenir Idex and ANR, France; DFG and AvH Foundation, Germany; Herakleitos, Thales and Aristeia programmes co-financed by EU-ESF and the Greek NSRF, Greece; BSF-NSF and MINERVA, Israel; NCN and NAWA, Poland; La Caixa Banking Foundation, CERCA Programme Generalitat de Catalunya and PROMETEO and GenT Programmes Generalitat Valenciana, Spain; Göran Gustafssons Stiftelse, Sweden; The Royal Society and Leverhulme Trust, United Kingdom.

In addition, individual members wish to acknowledge support from Armenia: Yerevan Physics Institute (FAPERJ); CERN: European Organization for Nuclear Research (CERN PIAS); Chile: Agencia Nacional de Investigación y Desarrollo (FONDECYT 1230812, FONDECYT 1230987, FONDECYT 1240864); China: Chinese Ministry of Science and Technology (MOST-2023YFA1605700), National Natural Science Foundation of China (NSFC - 12175119, NSFC 12275265, NSFC-12075060); Czech Republic: Czech Science Foundation (GACR - 24-11373S), Ministry of Education Youth and Sports (FORTE CZ.02.01.01/00/22_008/0004632), PRIMUS Research Programme (PRIMUS/21/SCI/017); EU: H2020 European Research Council (ERC - 101002463); European Union: European Research Council (ERC - 948254, ERC 101089007), Horizon 2020 Framework Programme (MUCCA - CHIST-ERA-19-XAI-00), European Union, Future Artificial Intelligence Research (FAIR-NextGenerationEU PE00000013), Italian Center for High Performance Computing, Big Data and Quantum Computing (ICSC, NextGenerationEU); France: Agence Nationale de la Recherche (ANR-20-CE31-0013, ANR-21-CE31-0013, ANR-21-CE31-0022), Investissements d'Avenir Labex (ANR-11-LABX-0012); Germany: Baden-Württemberg Stiftung (BW Stiftung-Postdoc Eliteprogramme), Deutsche Forschungsgemeinschaft (DFG - 469666862, DFG - CR 312/5-2); Italy: Istituto Nazionale di Fisica Nucleare (ICSC, NextGenerationEU); Japan: Japan Society for the Promotion of Science (JSPS KAKENHI JP22H01227, JSPS KAKENHI JP22H04944, JSPS KAKENHI JP22KK0227, JSPS KAKENHI JP23KK0245); Netherlands: Netherlands Organisation for Scientific Research (NWO Veni 2020 - VI.Veni.202.179); Norway: Research Council of Norway (RCN-314472); Poland: Ministry of Science and Higher Education (IDUB AGH, POB8, D4 no 9722), Polish National Agency for Academic Exchange (PPN/PPO/2020/1/00002/U/00001), Polish National Science Centre (NCN 2021/42/E/ST2/00350, NCN OPUS nr 2022/47/B/ST2/03059, NCN UMO-2019/34/E/ST2/00393, UMO-2020/37/B/ST2/01043, UMO-2021/40/C/ST2/00187, UMO-2022/47/O/ST2/00148, UMO-2023/49/B/ST2/04085, UMO-2023/51/B/ST2/00920); Slovenia: Slovenian Research Agency (ARIS grant J1-3010); Spain: Generalitat Valenciana (Artemisa, FEDER, ID-IFEDER/2018/048), Ministry of Science and Innovation (MCIN & NextGenEU PCI2022-135018-2, MICIN & FEDER PID2021-125273NB, RYC2019-028510-I, RYC2020-030254-I, RYC2021-031273-I, RYC2022-038164-I); Sweden: Swedish Research Council (Swedish Research Council 2023-04654, VR 2018-00482, VR 2022-03845, VR 2022-04683, VR 2023-03403, VR grant 2021-03651), Knut and Alice Wallenberg Foundation (KAW 2018.0458, KAW 2019.0447, KAW 2022.0358); Switzerland: Swiss National Science Foundation (SNSF - PCEFP2_194658); United Kingdom: Leverhulme Trust (Leverhulme

561 Trust RPG-2020-004), Royal Society (NIF-R1-231091); United States of America: U.S. Department of
562 Energy (ECA DE-AC02-76SF00515), Neubauer Family Foundation.

References

- [1] ATLAS Collaboration, *Performance of the ATLAS trigger system in 2015*, *Eur. Phys. J. C* **77** (2017) 317, arXiv: 1611.09661 [hep-ex] (cit. on pp. 3, 4).
- [2] ATLAS Collaboration, *Performance of electron and photon triggers in ATLAS during LHC Run 2*, *Eur. Phys. J. C* **80** (2020) 47, arXiv: 1909.00761 [hep-ex] (cit. on pp. 3, 7, 9, 18, 19).
- [3] ATLAS Collaboration, *Search for pairs of scalar leptoquarks decaying into quarks and electrons or muons in $\sqrt{s} = 13$ TeV pp collisions with the ATLAS detector*, *JHEP* **10** (2020) 112, arXiv: 2006.05872 [hep-ex] (cit. on p. 3).
- [4] ATLAS Collaboration, *Search for light neutral particles decaying promptly into collimated pairs of electrons or muons in pp collisions at $\sqrt{s} = 13$ TeV with the ATLAS detector*, *Eur. Phys. J. C* **85** (2025) 335, arXiv: 2407.09168 [hep-ex] (cit. on p. 3).
- [5] ATLAS Collaboration, *Observation of a new particle in the search for the Standard Model Higgs boson with the ATLAS detector at the LHC*, *Phys. Lett. B* **716** (2012) 1, arXiv: 1207.7214 [hep-ex] (cit. on p. 3).
- [6] ATLAS Collaboration, *Measurement of the Higgs boson mass in the $H \rightarrow ZZ^* \rightarrow 4\ell$ and $H \rightarrow \gamma\gamma$ channels with $\sqrt{s} = 13$ TeV pp collisions using the ATLAS detector*, *Phys. Lett. B* **784** (2018) 345, arXiv: 1806.00242 [hep-ex] (cit. on p. 3).
- [7] ATLAS Collaboration, *The ATLAS Experiment at the CERN Large Hadron Collider*, *JINST* **3** (2008) S08003 (cit. on pp. 3, 4).
- [8] ATLAS Collaboration, *ATLAS Insertable B-Layer: Technical Design Report*, ATLAS-TDR-19; CERN-LHCC-2010-013, 2010, URL: <https://cds.cern.ch/record/1291633> (cit. on p. 3), Addendum: ATLAS-TDR-19-ADD-1; CERN-LHCC-2012-009, 2012, URL: <https://cds.cern.ch/record/1451888>.
- [9] B. Abbott et al., *Production and integration of the ATLAS Insertable B-Layer*, *JINST* **13** (2018) T05008, arXiv: 1803.00844 [physics.ins-det] (cit. on p. 3).
- [10] ATLAS Collaboration, *Electron and photon performance measurements with the ATLAS detector using the 2015-2017 LHC proton–proton collision data*, (2019), arXiv: 1908.00005 [hep-ex] (cit. on p. 5).
- [11] ATLAS Collaboration, *Topological cell clustering in the ATLAS calorimeters and its performance in LHC Run 1*, *Eur. Phys. J. C* **77** (2017) 490, arXiv: 1603.02934 [hep-ex] (cit. on p. 5).
- [12] ATLAS Collaboration, *Improved electron reconstruction in ATLAS using the Gaussian Sum Filter-based model for bremsstrahlung*, ATLAS-CONF-2012-047, 2012, URL: <https://cds.cern.ch/record/1449796> (cit. on p. 5).
- [13] ATLAS Collaboration, *Electron reconstruction and identification in the ATLAS experiment using the 2015 and 2016 LHC proton–proton collision data at $\sqrt{s} = 13$ TeV*, *Eur. Phys. J. C* **79** (2019) 639, arXiv: 1902.04655 [physics.ins-det] (cit. on pp. 6, 11, 22).
- [14] J. Seixas, L. Caloba, M. Souza, A. Braga and A. Rodrigues, *Neural second-level trigger system based on calorimetry*, *Comput. Phys. Commun.* **95** (1996) 143, ISSN: 0010-4655 (cit. on p. 8).

- [15] ATLAS Collaboration, *ATLAS data quality operations and performance for 2015–2018 data-taking*, *JINST* **15** (2020) P04003, arXiv: [1911.04632](https://arxiv.org/abs/1911.04632) [[physics.ins-det](https://arxiv.org/archive/physics)] (cit. on p. 10).
- [16] S. Haykin, *Neural Networks and Learning Machines*, 3rd ed., Upper Saddle River, NJ, USA: Prentice Hall, 2008 (cit. on pp. 10, 11).
- [17] D. Nguyen and B. Widrow, ‘Improving the Learning Speed of 2-Layer Neural Networks by Choosing Initial Values of the Adaptive Weights’, *Proceeding of the International Joint Conference on Neural Networks*, vol. 3, San Diego, USA, 1990 21 (cit. on p. 10).
- [18] M. Riedmiller and H. Braun, ‘A Direct Adaptive Method for Faster Backpropagation Learning: The RPROP Algorithm’, *Proceedings of the IEEE International Conference on Neural Networks*, San Francisco, USA, 1993 586 (cit. on p. 10).
- [19] I. Goodfellow, Y. Bengio and A. Courville, *Deep learning*, 2016 800, URL: <http://www.deeplearningbook.org> (cit. on p. 11).
- [20] M. Aaboud et al., *Electron reconstruction and identification in the ATLAS experiment using the 2015 and 2016 LHC proton–proton collision data at TeV*, *The European Physical Journal C* **79** (2019) 1 (cit. on p. 11).
- [21] D. P. Kingma and J. Ba, *Adam: A method for stochastic optimization*, arXiv preprint arXiv:1412.6980 (2014) (cit. on p. 12).
- [22] ATLAS Collaboration, *Trigger monitoring and rate predictions using Enhanced Bias data from the ATLAS Detector at the LHC*, ATL-DAQ-PUB-2016-002, 2016, URL: <https://cds.cern.ch/record/2223498> (cit. on p. 15).
- [23] N. D. Gagunashvili, *Comparison of weighted and unweighted histograms*, ArXiv Physics e-prints (2006), arXiv:physics/0605123 (cit. on p. 22).
- [24] K. Pearson, *On the probability that two independent distributions of frequency are really samples from the same population*, *Biometrika* **8** (1911) 250 (cit. on p. 22).
- [25] T. D. Wickens, *Multiway contingency tables analysis for the social sciences*, Psychology Press, 2014 (cit. on p. 22).
- [26] C. Ma, G. Shan and S. Liu, *Homogeneity test for correlated binary data*, *PloS one* **10** (2015) e0124337 (cit. on p. 22).
- [27] ATLAS Collaboration, *ATLAS Computing Acknowledgements*, ATL-SOFT-PUB-2023-001, 2023, URL: <https://cds.cern.ch/record/2869272> (cit. on p. 24).